Evaluating and Improving Static Analysis Tools Via Differential Mutation Analysis

Abstract—**Many programming languages offer multiple static analysis tools that offer to detect faults in code without executing it. Understanding the strengths and weaknesses of tools, and performing direct comparisons of their effectiveness is difficult; it usually involves either manual examination of differing warnings on real code, or the bias-prone construction of artificial test cases. This paper proposes a novel automated approach to comparing static analysis tools, based on producing *mutants* of real code, and comparing mutation detection rates for tools to their warning rates on the original code. In addition to making tool differences quantitatively observable without extensive manual effort, this approach offers a new way to detect and fix omissions in a static analysis tool's set of detectors. We present an extensive comparison of three well-known Solidity smart contract static analysis tools, and show how using an automatic prioritization of our results allowed us to add three effective new detectors to the best of these. We also evaluate popular Java and Python static analysis tools and discuss their strengths and weaknesses.**

I. INTRODUCTION

Static analysis of code is one of the most effective ways to avoid defects in software, and, when security is a concern, is essential. Static analysis can find problems that are extremely hard to detect by testing, when the inputs triggering a bug are hard to find. Static analysis is also often more efficient than testing; a bug that takes a fuzzer days to find may be immediately identified. Users of static analysis tools often wonder which of multiple tools available for a language are most effective, and how much tools overlap in their results. Tools often find substantially different bugs, making it important to use multiple tools [32]. However, given the high cost of examing results, if a tool provides only marginal novelty, it may not be worth using, especially if it has a high false-positive rate. Developers of static analysis tools also want to be able to compare their tools to other tools, in order to see what detection patterns or precision/soundness trade-offs they might want to imitate. Unfortunately, comparing static analysis tools in these ways is hard, and would seem to require vast manual effort to inspect findings and determine ground truth on a scale that would provide statistical confidence.

Differential testing [41], [29], [56] is a popular approach to comparing multiple software systems offering similar functionality, but the wide divergence of possible trade-offs, analysis focuses, and the prevalence of false positives in almost all analysis results makes naïve differential testing not applicable to static analysis tools [14]. Mutation analysis[34], [15], [7] uses small syntactic changes to a program to introduce synthetic "faults," under the assumption that if the original version of a program is mostly correct, such changes will often

introduce a fault. For the most part, mutation analysis has been used to evaluate test suites by computing a mutation score, the fraction of mutants the suite detects, or "kills". Groce et al. [24], [25] proposed examining individual mutants that survive a rigorous testing and verification effort to detect and correct weaknesses in testing, and found bugs in a heavily-tested module of the Linux kernel [2] and a widely used Python file system. Recently, mutation analysis has been adopted in industrial settings, though not for actual examination of all surviving mutants [45], [33].

Combining a differential approach and mutation analysis offers a novel way to compare static analysis tools, one useful to users wishing to select a good tool or set of tools, to researchers interested in the impact of precision/soundness trade-offs or different intermediate languages, and to developers of static analysis tools hoping to improve their tools.

We can say that a static analysis tool kills a mutant when the *number of warnings or errors*, which we call *findings*, *increases with mutation*. In order to make this definition useful, we ignore informational or optimization related warnings (e.g., if a mutant is merely *stylistically* suboptimal this is not "finding a fault"). That is, a mutant is killed when a tool "finds more (unique) bugs" for the mutated code than for the unmutated code. This difference may be most easily interpreted when the original code produces no findings; we call such code *clean* (by analogy with Chekam et al.'s notion [10]). For non-clean code, a tool conceivably could detect the mutant, but only change a previously generated finding, not add an additional finding. However, even for non-clean code, *most detected mutants should produce a new warning*. We count findings, rather than consider their location or type, because some mutants cause a fault at a far-removed location. Forcing tools to produce an *additonal* warning is a conservative and automatable estimate of mutant detection.

The value of the differential comparison lies in a few key points. First, this is a measure that does not reward a tool that produces too many false positives. The tool cannot simply flag all code as having a problem or it will perform poorly at the task of *distinguishing* the mutated code from non-mutated, and presumably at least *more* correct, code. Based on verification and testing uses of mutation, it is safe to say that usually at minimum 40%, often 50-60%, and frequently up to 80%+ [2], [25], [39], of mutants are not semantically equivalent to the original code [42], [31], [49], so the task presented to a static analysis tool is simply the core functionality of static analysis: *to distinguish faulty from correct code without execution.* Obviously, many faults cannot be identified statically without

Fig. 1: Mutants killed by three static analysis tools.

```
contract SimpleStorage {
    uint storedData;
    function set(uint x) public {
        storedData = x;
    }
    function get() public view returns (uint) {
        return storedData;
    }
}
```

Fig. 2: A simple example Solidity smart contract

a complete specification, or without unreasonable analysis cost and precision, but the measure of performance here is *relative* to other tools applied to the same code; this is primarily a *differential* approach. While many mutants cannot be detected statically, the ones that *are* tend to be *true positives*: *if* they were real code changes, they would be faults. *We manually confirmed that for a large portion of the detected mutants in our experiments, the changes were indeed ones that would be real faults if present in the code.*

Second, and critically, this is an *automatable* method that can provide an evaluation of static analysis tools over a large number of target source code files, without requiring human effort to classify results as real bugs or false positives. It is not clear that any other fully automatic method is competitively meaningful; it is possible that methods based on code changes from version control provide some of the same benefits, but these require classification of changes into bug-fixes and non-bug-fixes, and of course require version control history. Also, history-based methods will be biased towards precisely those faults humans or tools already in use were able to detect and fix. Rather than the hundreds [35] or at most few thousand of faults [53] in benchmark defect sets, our approach enables the use of *many tens of thousands* of hypothetical faults.

It is the combination of differential comparison and mutation that is key. Differential comparison of tools, as noted above, is not really meaningful, without additional effort; naïve methods simply will not work [14]. Consider a comparison of the number of findings between two tools over a single program, or over a large set of programs. If one tool emits more warnings and errors than another, it may mean that the tool is more effective at finding bugs; but it may also mean that it has a higher false positive rate. Without human examination of the individual findings, it is impossible to be sure. Using mutants, however, provides a foreground to compare to this background. In particular, for a large set of programs, the most informative result will be when 1) tool A reports fewer findings on average than tool B over the un-mutated programs but 2) tool A also detects more mutants. This is strong evidence that A is simply better all-around than B; it likely has a lower false positive rate *and* a lower false negative rate, since it is hard to construct another plausible explanation for reporting *fewer* findings on un-mutated code while still detecting *more* mutants. Our method (see the *mutant ratio* defined below) provides a quantitative measure of this insight.

Finally, even when tools have similar quantitative results, examining individual mutants killed by one tool but not by

another allows us to understand strengths and weaknesses of the tools, in a helpful context: the difference between the un-mutated code and mutated code will always be small and simple. Moreover, simply looking at how much two tools agree on mutants can answer the question: given that I am using tool A, would adding tool B be likely be worthwhile? Interested users, e.g. security analysts, can inspect the differences to get an idea of the particular cases when a tool might be most effective, but a more typical user can simply look at a Venn diagram of kills like that shown in Figure 1. Consider hypothetical tools A, B, and C. A and B produce similar numbers of findings on the code in question, while tool C produces an order of magnitude more findings. Tool A is likely the most important tool to make use of; it detected more mutants than any other tool, and more than twice as many mutants were killed by A alone than by B alone. However, also running tool B is well justifed. B does not do as well as A, but it is the only tool that detects a large number of mutants, and most mutants it detects are unique to it. Finally, Tool C *may* not be worth running, since its poor performance on mutants but high finding rate suggests it may be prone to missing bugs and to false positives. It might be a good idea to just look at the 27 mutants detected by C alone: *if* they represent an important class of potential problems (perhaps C specialized in detecting potentially non-terminating loops), then C might be useful, but if the first few mutants inspected are false positives, then C is likely not useful.

More concretely, consider the code in Figure 2 [1]. The Universal Mutator tool [27], [28], which has been extensively tuned for Solidity's grammar (though not to target any particular vulnerabilities), and is the only smart contract mutation tool referenced in the Solidity documentation (https://solidity.readthedocs.io/en/v0.5.12/resources.html), produces seven valid, non-redundant (by Trivial Compiler Equivalence [42]) mutants for it. Both the public version of Slither [19] and SmartCheck [51] (two popular smart contract static analysis tools) produce a small number (three and two, respectively) of low-severity, informational, warnings for this code. Both tools also detect four of the seven mutants (here the number of warnings increases, and the additional warnings are clearly driven by the mutation change) . However, only one of the mutants detected is common to both tools: both tools detect changing the `return` statement in the `get` function to a call to `selfdestruct` the smart contract, deleting it. Slither, but not SmartCheck, also detects replacing the assignment of `storedData` in `set` with either a `selfdestruct` or

`revert`, or simply removing it altogether. SmartCheck, on the other hand, detects removing the `return` in `get` or replacing it with a `revert`, or removing the `public` visibility modifier for `get` [1]. If we restrict our analysis to findings with a severity greater than informational, SmartCheck detects no mutants of the contract, while Slither still reports that some mutants *allow an arbitrary caller to cause the contract to destroy itself*. Given that both tools, ignoring informational results, detect no problems with the original code, and only Slither detects any problems with the mutants, we can say that Slither performs better for this contract.

Comparing mutant results also leads to the idea of *improving* static analysis tools by examining mutants detected by another tool, and thus known to be in-principle detectable. Improving tools by adding detectors is useful because, even if all tools had the same set of detectors, they would not all report the same bugs; different choices in intermediate language and tradeoffs made to avoid false positives may make the use of multiple tools with similar detectors essential for thorough analysis. And if one tool simply has a superior engine, it is beneficial to users that the "best" tool incorporate *all* detection rules. However, as with efforts to improve test suites, manually searching through all mutants can be an onerous task, especially for large-scale evaluations. We therefore introduce the idea of *prioritizing* mutants to make it easier to inspect *different* weaknesses in tools.

A general objection to our approach is that mutants may differ substantially from "real" faults, in some way. This is certainly true, in a sense [23], but for static analysis purposes we believe it does not matter. The real risk is that some mutation operators align with patterns a particular tool identifies, biasing the evaluation in favor of that tool. Such faults may be dis-proportionately present in mutants vs. real code. However, we consider this unlikely. The vast majority of applied mutation operations for all of our experiments were highly generic, and do not plausibly represent a pattern in which some tool might specialize. Code deletions, the most common kind of mutation by far, leave no "trace" for a tool to match against, but only an omission, so cannot be subject to this concern. Changing arithmetic and comparison operators and numeric constants (incrementing, decrementing, or changing to 0 or 1) account for most of the non-statement-deletion mutants, and it is difficult to imagine how any tool could unfairly identify these.

This paper offers the following contributions:

- We propose a differential approach to comparing static analysis tools based on the insight that program mutants are easy-to-understand, likely-faulty, program changes.
- We propose a definition of mutant killing for static analysis.
- We introduce a simple scheme for prioritizing mutants that helps users understand and use the results of analysis.
- We apply our method to an extensive, in-depth comparison of three Solidity smart contract analysis tools, and show

how prioritization allowed us to easily identify (and build) three new detectors for the most effective of these tools.
- We also provide results for popular Java and Python static analysis tools, further demonstrating our approach and showing strengths and weaknesses of these tools.

While there are limitations to using differential mutation analysis to compare/improve static analysis tools, it scales to basing comparisons on many real software source files and very many "faults," but still offers some of the advantages of having humans establish ground truth.

II. DIFFERENTIAL MUTATION ANALYSIS

The proposed approach is simple in outline:

1) Run each tool on the unmutated source code target(s), and determine the *baseline*: the number of (non-informational/stylistic) findings produced.
2) Generate mutants of the source code and run each tool on each mutant. Consider a mutant killed if the number of findings for the mutated code is greater than the number for the baseline, un-mutated code.
3) Compute, for each tool, the *mutant ratio*: the mutation score ($\frac{|killed|}{|mutants|}$) divided by (mean) baseline. If it is zero, use a baseline equal to either one or the lowest non-zero baseline for any tool in the comparison set[2].
4) (Optional): Discard all mutants not killed by at least one tool and all mutants killed by all tools. What remains allows *differential* analysis. Examine the remaining mutants in the difference in *prioritized* order.

The most important step here is the computation of the *mutant ratio*, which tells us about the ability of a tool to produce findings *for mutants*, relative to its tendency to produce findings in general. If a tool has a tendency to produce large numbers of findings compared to other tools, and this is paired with a tendency to detect more mutants as well, then the tool will not be penalized for producing many findings. Assuming that real faults are relatively rare in the original, un-mutated code, the best result and best (highest) mutant ratio will be for a tool that produces comparatively few findings for un-mutated code, but detects a larger portion of mutants than other tools; the worst result will be a tool that produces lots of findings, but detects few mutants. We will actually see some examples of this worst case.

A. Prioritizing Mutants

One goal of our approach is to make it easy for tool developers to examine cases where one tool kills a mutant and another fails to, in order to identify patterns for new detectors or analysis algorithm problems. Dedicated developers may also simply want to scan all mutants their tool does not kill, for the same purpose, analogous to what Groce et al. have proposed for automated verification and testing [24], [25]. Security analysts and other expert users who are not developers may also wish to do this, to better understand tool strengths and weaknesses.

Unfortunately the full list of unkilled mutants, or differentially unkilled mutants, is likely to be both large and highly redundant. In our results below, only one of 9 tools we examined killed fewer than 1,000 mutants it alone detected. Any cross-tool comparison is thus likely to involve hundreds or thousands of mutants.

The problem of identifying unique "faults" (tool weaknesses) in this situation is very similar to the *fuzzer taming* problem in software testing, as defined by Chen et al. [11]: "Given a potentially large collection of test cases, each of which triggers a bug, rank them in such a way that test cases triggering distinct bugs are early in the list." [11]. Their solution was to use Gonzalez' Furthest-Point-First [21] (FPF) algorithm to *rank* test cases so that users can examine very different test cases as quickly as possible. An FPF ranking requires a distance metric d, and ranks items so that dissimilar ones appear earlier. The hypothesis of Chen et al. was that dissimilar tests, by a well-chosen metric, will also fail due to different faults. FPF is a greedy algorithm that proceeds by repeatedly adding the item with the *maximum minimum distance to all previously ranked items*. Given an initial seed item r_0, a set S of items to rank, and a distance metric d, FPF computes r_i as $s \in S : \forall s' \in S : min_{j<i}(d(s, r_j)) \geq min_{j<i}(d(s', r_j))$. The condition on s is obviously true when $s = s'$, or when $s' = r_j$ for some $j < i$; the other cases for s' force selection of *some* max-min-distance s.

In order to apply FPF ranking to examining mutants, we implemented a simple, somewhat *ad hoc* distance metric and FPF ranker. Our metric d is the sum of a set of measurements. First, it adds a similarity ratio based on Levenshtein distance [40] for (1) the *changes* (Levenshtein edits) from the original source code elements to the two mutants, (2) the two original source code elements changed (in general, lines), and (3) the actual output mutant code. These are weighted with multipliers of 5.0, 0.1, and 0.1, respectively; the type of change (mutation operator, roughly) dominates this part of the distance, because it best describes "what the mutant did"; however, because many mutants will have the same change (e.g., changing + to −, the other ratios also often matter. Our metric also incorporates a measure of the distance in the source code between the locations of two mutants. If the mutants are to different files, this adds 0.5; it also adds 0.025 times the number of source lines separating the two mutants if they are in the same file, but caps the amount added at 0.25.

We do not claim this is an optimal, or even tuned, metric. Devising a better metric is left as future work, we only wish to show that even a hastily-devised and somewhat arbitrary metric provides considerable advantage over wading through an un-ordered list of mutants, and introduce the idea of using FPF for mutants, not just for tests: FPF is useful for failures in general, however discovered.

III. EXPERIMENTAL RESULTS

Our primary experimental results are a set of comparisons of tools using our method, for three languages: Solidity (the most popular language for smart contracts), Java, and Python.

We used the Universal Mutator tool for all experiments; for Solidity and Python, we believe the Universal Mutator is simply the best available tool. For Java, PIT [12] is more popular, but does not produce source-level mutants, needed for PMD and for manual inspection of results. Universal Mutator includes a large set of mutation operators, some unconventional (e.g., swapping order of function arguments) but based on real-world bugs; the complete set is described by regular expressions at https://github.com/agroce/universalmutator/tree/master/universalmutator/static. However, most mutants that were detected came from a small set of commonly-used operators [3], particularly 1) code deletion and 2) operator, conditional, and constant replacements.

We used our results to answer a set of research questions:

- **RQ1:** Does mutation analysis of static analysis tools produce actionable results? That is, do raw mutation kills serve to distinguish tools from each other, or are all tools similar?
- **RQ2:** Does our approach provide additional information beyond simply counting findings for the original, unmutated analyzed code? Do *ratios* differ between tools?
- **RQ3:** Do the rankings that raw kills and ratios establish agree with other sources of information about the effectiveness of the evaluated tools?
- **RQ4:** Do tools detect more mutants in programs for which they produce no warnings, initially?
- **RQ5:** Are mutants distinguishing tools usually flagged due to real faults, where the finding is related to the introduced fault; that is, are our results usually *meaningful*?
- **RQ6:** Do individual mutants, prioritized for ease of examination, allow us to identify classes of faults that different tools are good at/bad at, and use this information to improve tools? How does this compare to using mutants that have *not* been prioritized?

In particular, we consider **RQ2** to be of critical importance; if the mutant ratios for tools differ, then this is clear evidence that our hypothesis that the tendency of mutants to be faults, and to expect that mutated code will, by a more precise and accurate tool, be flagged as problematic more often than non-mutated code, holds. This expectation that (some subset of the) mutants can serve as proxies for real, detectable faults is the core concept of our approach. **RQ4** addresses a concern briefly mentioned in the introduction: it is possible that warnings for the original code interfere with our definition of detection. The ideal case for our approach is when a tool reports no findings for un-mutated code, and reports a finding when the mutant is introduced. Chekam et al. showed that the "clean program assumption" for testing is a threat to the validity of investigations of the relationship between coverage and fault detection [10], but we show that this is unlikely to be the case for our approach. Our answer to **RQ5** is somewhat inherently qualitative and incomplete; we cannot analyze all mutants on which results are based manually, and understanding the mutants and tool warnings completely would require deep understanding of all the subject programs. However, in many cases, the impact of a mutant is clear, and the reason for

warnings is obvious. This was often enough the case that, as we discuss below, we are confident mutants that distinguish tools are meaningful (missed) opportunities for static analysis tools. For **RQ6** we have only a preliminary answer.

A. Solidity Smart Contract Tools

1) Smart Contracts and Smart Contract Static Analysis: Smart contracts are autonomous code instruments, usually operating on a blockchain, that often have critical responsibilities such as facilitating and verifying (large) financial services transactions, tracking high-value physical goods or intellectual property, or even controlling "decentralized organizations" with multifarious aspects. Security and correctness are thus critical in the smart contract domain, and static analysis is a key way to ensure allocation of high-value resources is not compromised. The most popular smart contract platform, by far, is the Ethereum blockchain, and the Solidity smart contract language [8], [55]; the Ethereum cryptocurrency has a market capitalization as we write of over $100 billion dollars, largely fueled by interest in the smart contract functionality. Ethereum contracts have been the targets of widely publicized attacks, with large financial consequences [50], [46]. A recent paper examining results from 23 professional security audits of Solidity contracts argues that effective static analysis is a major key to avoiding such disasters in the future [26].

2) Static Analysis Tools Compared: We analyzed three well-known tools for static analysis of Solidity smart contracts: Slither [19], SmartCheck [51], and Securify [54]. Slither, based on an SSA-based intermediate language (SlithIR [19]) is an open-source tool from Trail of Bits. SmartCheck, developed by SmartDec, translates Solidity source directly to an XML-based representation, then uses *XPath* patterns to define problems. Securify, from SRI Systems Lab at ETH Zurich, works at the bytecode level, first parsing and decompiling contracts, then translating to *semantic facts* in order to look for problems.

3) Smart Contract Selection: We could have used a set of high-transaction contracts, or known-important contracts to validate our approach. However, we knew that one of our goals in the Solidity experiments was to actually improve a mutation analysis tool, and the developers of the static analysis tools use exactly such benchmarks to validate their tools. Basing our improvements on mutants of the contracts used for evaluation of proposed detectors would introduce a serious bias in our favor: we would be more likely to produce detectors that would have true positives and few false positives on the benchmark contracts. We therefore instead selected 100 random contracts for which EtherScan (https://etherscan.io/) has source code, and used this (quite arbitrary) set of contracts from the actual blockchain to compare tools and identify opportunities for improvement. The collected contracts had a total of 15,980 non-comment source lines, as measured by `cloc`, with a mean size of 159.8 LOC and a median size of 108 LOC. The largest single contract had 1,127 lines of code. The Universal Mutator generated 46,769 valid mutants for these 100 contracts.

4) Analysis Results: Figure 3 shows the mutants killed by the Solidity analysis tools. Table I provides numeric details


Fig. 3: Mutants killed by Solidity static analysis tools.

Tool	Findings		Mutation Score		Mutant Ratio
	Mean	Median	Mean	Median	
Slither	2.37	1.0	0.09	0.09	0.038
Clean (39)	-	-	0.11	0.11	-
SmartCheck	1.89	1.0	0.05	0.05	0.026
Clean (27)	-	-	0.03	0.01	-
Securify	24.65	17.0	0.03	0.02	0.001
Clean (5)	-	-	0.00	0.00	-

TABLE I: Solidity tool results over all contracts.

of the results, including the *ratio* for each tool, adjusting its mutation scores by its' general tendency to produce findings. The second row for each tool shows the number of contracts for which it reported no findings, and the mutation scores over those contracts, only. A user examining these results would suspect that Slither and SmartCheck are both useful tools, and should likely both be applied in a high-risk security-sensitive context like smart contract development. A user might also suspect that the large number of findings produced, and smaller number of mutants killed, for Securify, means that whether to apply Securify is a more difficult decision. On the one hand, Securify does detect nearly as many mutants it alone can identify as SmartCheck. The large number of findings, and very bad mutant ratio, however, lead us to suspect that many of these "detected" mutants are false positives (or, at least, that the problem is not the one Securify identifies). Extracting the signal from Securify's noise will be difficult. We also note that while running Slither and SmartCheck on all 46,769 valid mutants was relatively quick (it took about 6 days sequential compute time for Slither and 3 days for SmartCheck, i.e., about 5-15 seconds per mutant for both tools), Securify often required many hours to analyze a mutant, and frequently required a few days to analyze a mutant; the full analysis required over three months of compute time. However, similar statistical results would almost certainly be produced by running as few as 1,000 mutants, randomly selected, due to the implications of Tchebysheff's inequality [22], a technique that should improve scalability for all such analyses.

If we consider results at the individual contract level, the overall picture is in some ways even clearer. Slither detected more mutants than Securify for 84 of the contracts, and more than SmartCheck for 85 of the contracts. SmartCheck detected more mutants than Securify for 83 of the contracts. Comparing ratios instead, Slither was better than Security for 97 contracts, but better than SmartCheck for only 75 of the contracts. SmartCheck's ratio was better than Securify's for 95 contracts.

The standard deviation of contract raw scores was 0.05 for Slither, 0.03 for SmartCheck, and 0.04 for Securify.

For our research questions, **RQ1** is clearly answered in the affirmative. Figure 3 shows that the tools address quite different problems, with all tools reporting far more uniquely detected mutants than mutants in common with other tools. There are only 18 mutants detected by all tools, all of them involving replacement of `msg.sender` (the caller of a smart contract, which may be another smart contract) with `tx.origin` (the original initiator of a sequence of blockchain calls, a "human" account). Use of `tx.origin` is often (though not always) a bad idea, and can lead to incorrect behavior, so it is not surprising tools all recognize some misuses of it.

RQ2 is also answered in the affirmative. Counting findings for un-mutated code might suggest that Securify is the best tool, by a wide margin, but in the context of its near-zero mutant ratio, we must suspect (and we partially manually confirmed) that many of the warnings are false positives. Slither has the best mutant ratio, but the margin between it and SmartCheck confirms that both tools likely provide value; we note there is a 25% chance that SmartCheck has a better ratio than Slither for an individual contract.

For **RQ3**, there are only a few tool comparisons in the literature; this is probably due to the fast-moving nature of the blockchain analysis world; the oldest of these tools' publication dates is 2018. The most extensive is that of Durieux et al. [18], though it unfortunately was unable to provide anything other than an implicit look at false positives, somewhat limiting its practicality. Slither detected 17% of known vulnerabilities in their analysis, vs. 11% for SmartCheck and 9% for Securify. Slither and SmartCheck were also among the four (out of 9) tools that detected vulnerabilities in the most categories; Securify was not. The overall recommendation of Durieux et al. was to use a combination of Slither and Mythril [13] for contract analysis. Parizi et al. [43] also offer a ranking of tools, and determined that SmartCheck was the most effective, and far more so than Securify; unfortunately, they did not include Slither in their set of evaluated tools.

The Slither paper [19] also provides an evaluation of all three tools. Their findings counts differ from ours because of different choices (we threw out merely informational results), but these are unrelated to mutation analysis, in any case. The evaluation only considered reentrancy faults [5], [26] (which are sometimes, but only rarely, introduced by mutants). For reentrancy, Slither performed best on two real-world large contracts, finding subtle bugs in both, SmartCheck detected the problem in one of the two, and Securify detected neither. For a set of 1,000 contracts, SmartCheck had a high false positive rate (over 70%) but detected more actual reentrancies (209) than Slither (99) or Securify (6). On the other hand, Slither's low false positive rate of 11% makes its results possibly more useful in practice.

For **RQ4**, on the changes seen when restricting analysis to clean contracts, Slither did slightly better at detecting mutants when the original contract was clean for Slither, and the other two tools did somewhat worse on contracts for which they reported no findings. For the three contracts clean for all tools, Slither performed almost exactly as it did over contracts in general, and the other tools performed worse, by about the same margin as they did for their own clean contracts. For our approach, we only need a weak version of the "clean program assumption": the threat is that kills may be under-reported for non-clean programs, due to interference with findings for the original code. It is not a problem if mutation scores are *worse* for programs where a tool reports no findings for the un-mutated code. We therefore, for smart contracts, find no threat to our approach arising from the presence of findings on un-mutated code. We speculate that "clean" results for some tools result from contracts where the tool has trouble with the contract code, but does not actually crash; Slither may do better on clean code because it has fewer such failures, and clean contracts are probably somewhat easier to analyze.

Following the method proposed in Section II, for **RQ5** we focused on examining mutants detected by at least one tool, but not detected by all tools, the only ones that actually influence the comparison of tools. The vast majority of the mutants in this set were meaningful semantic changes a static analysis tool could be expected to detect, and the findings produced by tools were relevant to the nature of the fault. We do not believe that *all* mutants represent definite faults; some are harmless but unusual code changes. Many cases where use of `tx.origin` in place of `msg.sender` was flagged seem to us to be strange, but not necessarily incorrect, code. On the other hand, it is not at all unreasonable for tools to report such notably strange code. Our estimate is that, ignoring `tx.origin` cases, *at least 70% of the mutants detected by one, but not all, tools, represent realistic bugs*, and failure to detect is roughly equally due to missing detectors and imprecise analysis.

Because our random contracts' quality might be low, we also checked our results on 30 contracts from the Solidity documentation, the *Mastering Ethereum* book [4], and a handful of selected, recent, higher quality blockchain contracts. Slither had a mean mutation score of 0.11, vs. 0.04 for SmartCheck and 0.01 for Securify. Associated mutant ratios were 0.38, 0.08, and 0.007. Mutant kill overlap was also similar; in fact, Figure 1 shows the results: Slither is A, SmartCheck is B, and Securify is C.

5) Improving Slither (**RQ6**): Based on the differential mutation analysis, we identified three low-hanging fruit to improve the performance of Slither. We chose Slither in part because it seems to have a better underlying intermediate language and analysis engine, and thus is likely to produce better results for the same rule than the other tools. The process was simple. First, we produced a list of all mutants killed by either SmartCheck or Securify, but not killed by Slither. We then applied the prioritization method based on the FPF algorithm and the distance metric described in Section II-A, and examined the mutants in rank order. Many of the mutants were difficult to identify as true or false positives, absent context. Some opportunities for enhancement were clear, but seemed likely to require considerable effort to implement

Mutant showing Boolean constant misuse.

```
if (!p.recipient.send(p.amount)) \{  // Make the payment
==>          if (true) \{  // Make the payment
if (true) \{  // Make the payment
```

Mutant showing Type-based tautologies.

```
require(nextDiscountTTMTokenId6 >= 361 \&\& ...);
==>  ...361...==>...0...
require(nextDiscountTTMTokenId6 >= 0 \&\& ...);
```

Mutant showing Loss of precision.

```
byte char = byte(bytes32(uint(x) * 2 ** (8 * j)));
==>  ...*...==>.../...
byte char = byte(bytes32(uint(x) * 2 ** (8 / j)));
```

Fig. 4: Examples of mutants leading to new detectors.

without producing a large number of false positives. For example, Securify often detected when an ERC20 token contract's guard preventing making the special 0x0 address the owner of a contract was removed, and issued the error `Violation for MissingInputValidation`. Detecting such missing guards is probably useful, but doing so without producing false positives is non-trivial. We wanted to show that mutants could identify *useful* but *easy to implement* missing detectors. Examining only a few mutants, we identified three:

1) **Boolean constant misuse:** This detector flags code like `if (true)` or `g(b || true)` (where g is a function that takes a Boolean input). Constant-valued conditionals tend to indicate debugging efforts that have persisted into production code, or other faults; there are almost no circumstances where a conditional should not vary with state or input. This detector is actually split into two detectors, one for this serious issue, and an informational/stylistic detector that flags code such as `if (x == true)`, which is merely difficult to read.

2) **Type-based tautologies:** A type-based tautology is again a case where a Boolean expression has a constant value, but this is not due to misuse of a Boolean constant, but is instead due to the *types* in a comparison. For example, if x is an unsigned integer type, the comparison `x >= 0` is always true and `x < 0` is always false. This detector is a generalization of the SmartCheck detector https://github.com/smartdec/smartcheck/blob/master/rule_descriptions/SOLIDITY_UINT_CANT_BE_NEGATIVE/, modified to actually compute the ranges of types and identify other cases such as `y < 512` where y's type is `int8`.

3) **Loss of precision:** Solidity only supports integer types, so performing division before multiplication can introduce avoidable rounding. This is a fairly important problem, given Solidity code often performs critical financial calculations. SmartCheck provides a detector for such precision losses https://github.com/smartdec/smartcheck/blob/master/rule_descriptions/SOLIDITY_DIV_MUL/.

All three of these detectors were submitted as PRs, vetted over an internal benchmark set of contracts used by the Slither developers to evaluate new detectors, and accepted for

release in the public version of Slither. All three detectors produce some true positives (actual problems, though not always exploitable) in benchmark contracts, have acceptably low false positive rates, and were deemed valuable enough to include as non-informational (medium severity) detectors. The first mutants in prioritized rank exhibiting the issues, shown above, were the 2nd, 9th, and 12th non-statement-deletion mutants ranked for SmartCheck, out of over 800 such mutants. Using our prioritization, it was possible to identify these issues by examining fewer than 20 unkilled mutants. Without prioritization, on average a developer would have to look at more than 200, 80, and 400 mutants, respectively, to find instances of these problems. Interestingly, the very fact that these instances are "needles in a haystack" among the mutants not killed by Slither means that the results in Figure 3 and Table I are almost unaltered by our improvements to Slither: our analysis is fairly robust to modest tool improvements, unless added detectors account for a large number of mutants not detected by the tool. Adding such detectors will also only improve mutation ratio if they do not add many false positives. Substantial changes in results therefore require adding very effective (for mutants) detectors that seldom trigger for correct code (or at least trigger much less than for mutants). "Cheating" with respect to a mutation benchmark is thus, we hope, very difficult.

There were 92 separately ranked statement deletion mutants also. These, however, could all be ignored, as they were almost entirely duplicates related to the missing-return statement detector. If this detector were not already present as a private Slither detector, it would also be a good candidate for addition to the tool. Our three submitted detectors were not present as private detectors, and only one (the type-based tautology detector) had even been identified, via a GitHub issue, as a potential improvement (and only in the private version of Slither). Combining statement deletion mutants with other mutants only moved the mutants we used down to 3rd, 11th, and 14th positions. By default we rank statement deletions separately, since such mutants are usually easier to understand and evaluate, and in testing (but not static analysis) they are likely to be the most critical faults not detected.

Examining the first 100 mutants in the unprioritized lists for SmartCheck and Securify, ordered by contract ID and mutant number (roughly source line mutated) we were unable to identify *any* obviously interesting mutants, suggesting that it is indeed hard to use mutation analysis results without prioritization. A large majority of the mutants we inspected involved either the missing `return` problem noted in the introduction, or replacing `msg.sender` with `tx.origin`; Slither has a detector for misuses of `tx.origin`. SmartCheck and Securify tend to identify most (though not all) uses of `tx.origin` as incorrect, while Slither has a more selective rule, intended to reduce false positives. It is hard to scale our efforts here to a larger experiment, since writing and submitting changes to static analysis tools is always going to be a fairly onerous task, but we believe that our successful addition of new detectors, and the ease of identifying candidate detectors using mutant prioritization supports a limited affirmative answer to **RQ6**.

Fig. 5: Mutants killed by Java static analysis tools.

| Tool | Findings | | Mutation Score | | Mutant Ratio |
	Mean	Median	Mean	Median	
SpotBugs	28.93	14.00	0.07	0.07	0.002
Clean (3)	-	-	0.05	0.06 -	
PMD	53.73	32.00	0.07	0.07	0.001
Clean (0)	-	-	-	-	-
Infer	11.60	3.00	0.00	0.00	0.000
Clean (6)	-	-	0.00	0.00	

TABLE II: Java tool results over all projects.

B. Java Tools

1) Static Analysis Tools Compared: For Java, we again compared three tools. SpotBugs (https://spotbugs.github.io/) is the "spiritual successor" of FindBugs [6], [47]. PMD (https://pmd.github.io/) [47] is an extensible cross-language static code analyzer. FaceBook's Infer (https://fbinfer.com/) [16] focuses on diff-related detection of serious errors (concurrency, memory safety, and information flow).

2) Project Selection: For Java and Python, we did not have to worry about invalidating tool improvements by basing our results on benchmark code. We therefore aimed to use realistic, important source code. We selected top GitHub projects (defined by number of stars) for each language, and removed projects with fewer than 5 developers or less than six months of commit history (as well as projects that did not build). For Java, we analyzed the top 15 projects satisfying our criteria, with a maximum of 623,355 LOC and a minimum of 3,957 LOC, and a total size of 1.8 million LOC. Because the Universal Mutator does not "know" Java syntax, and Java is very verbose, the Java compiler rejected a large number of the generated mutants (e.g., deleting declarations). We still, due to the huge size of the source files and thus number of mutants (and time to compile full projects), restricted our analysis to files where Universal Mutator's implementation of TCE [42] for Java was useful, i.e. individual files that could be compiled and the bytecode compared, leaving us with just over 70,000 mutants, ranging from 136 to 10,016 per project.

3) Analysis Results: Figure 5 shows the mutants killed by the Java analysis tools, and Table II provides numeric results for projects and clean projects, respectively. At the individual project level, Infer was never best; PMD had a better raw score than SpotBugs for 10/15 projects, but SpotBug had a better *ratio* for 11. There is likely a tradeoff between verbosity and precision. Standard deviation in project scores was 0.05 for SpotBugs, 0.04 for PMD, and 0 for Infer.

In terms of **RQ1**, the raw kills results suggest there is considereable value in running both SpotBugs and PMD. Both

Fig. 6: Discussions of Java static analysis tools.

produce a large number of unique detections, though PMD produces about 20% more than SpotBugs. Infer on the other hand, is only able to detect two mutants, but these are unique; both were, however, spurious concurrency warnings. It may be that the diff sizes in our code were simply too small for Infer's approach. **RQ2** is also answered in the affirmative. While the raw kills for SpotBugs are not as good as for PMD, it had fewer findings, giving it a mutant ratio approximately twice that of PMD.

We note that SpotBugs crashed for many more Java programs than PMD and Infer (neither crashed for any original file in our experiments). SpotBugs "failed to detect" 23,000 of the mutants because it did not process the un-mutated file for 383 files, over 12 of the 15 projects—just over 23% of the 1,664 total files. Removing files for which SpotBugs failed, however, did not dramatically change results; SpotBugs' mean mutation score rose to 0.09, and mutant ratio rose to 0.003, but PMD's mean score rose to 0.10, and mutant ratio to 0.002.

For **RQ3**, to our knowledge there is no academic comparison of all three tools; one study dates from 2004 [47], used FindBugs, not SpotBugs, and reached no strong conclusions with respect to FindBugs vs. PMD; a more recent study found that SpotBugs outperformed Infer for Defects4J [35] bugs [32], but did not compare to PMD. However, the user postings listed in Figure 6, plus personal communications with security analysts who use these tools [20] supported some basic conclusions. SpotBugs is perhaps the best tool for finding bugs; PMD focuses more on stylistic issues and has a weaker semantic model. Running both is definitely recommended, as neither is extremely effective. Infer is closer to a model checker focusing on resource leaks than a truly general-purpose tool, arguably. In fact, we suspect Infer would perform much better if we used mutation operators targeting some important subtle Java bugs; Infer is probably not a good *general-purpose* static analysis tool for Java. Note that we did not use Infer's experimental detectors. For **RQ4** there were very few clean projects, but we see no evidence that non-clean code is a source of degradation in mutation detection.

For Java, again, the large majority ($> 75\%$) of randomly chosen mutants in the tool difference set we inspected for **RQ5** were definitely meaningful, essentially "real faults." In particular, for Java, the large majority of mutants involved either deleted method calls or changes to conditionals (e.g., `== null` to `!= null`) that would clearly introduce potential null pointer exceptions (NPEs), and such a possible NPE was the produced finding. While the differences between Solidity tools were often due to different detectors, the Java differences seemed mostly rooted in analysis engine methods; all tools aim

Fig. 7: Mutants killed by Python static analysis tools.

Tool	Findings		Mutation Score		Mutant
	Mean	Median	Mean	Median	Ratio
Pylint	10.63	5.50	0.31	0.34	0.03
Clean (4)	-	-	0.49	0.47	-
pyflakes	1.46	0.00	0.14	0.12	0.09
Clean (14)	-	-	0.14	0.12	
PyChecker	11.6	0.00	0.01	0.00	0.00
Clean (4)	-	-	0.00	0.00	-

TABLE III: Python tool results over all projects.

to warn about potential NPEs. Because the number of mutants we could examine and understand was smaller, we are less confident in making a probabilistic estimate than with Solidity, but it was clear the basis of the comparison was primarily realistic faults, which some tools detected and others did not.

C. Python Tools

1) Static Analysis Tools Compared: We compared three widely used and well-known Python tools: Pylint https://www.pylint.org/ (probably the most widely used of Python bug finding tools), pyflakes https://pypi.org/project/pyflakes/, designed to be faster, lighter-weight, and more focused on bugs (without configuration) than Pylint, and PyChecker http://pychecker.sourceforge.net/, an older, but still used tool.

2) Project Selection: For Python, we analyzed the top 25 GitHub projects by our criteria (see above), due to the smaller size of Python projects. These ranged in size from 137 LOC to 29,339 LOC, with a total size of about 75 KLOC, and a mean size of 3,185 LOC. We analyzed 158,418 valid, non-TCE redundant, mutants taken from these programs.

3) Analysis Results: Figure 7 shows the mutants killed by the Python analysis tools, and Table III provides numeric results for projects and clean projects, respectively. At the project level, Pylint was better than pyflakes and PyChecker for 21 and 24 projects, respectively, by raw score; pyflakes was better than Pylint and PyChecker for 3 and 24 projects, respectively; PyChecker was never better than another tool by raw score. Switching to ratio measures, Pylint was better than pyflakes and PyChecker for 11 and 23 projects, respectively; pyflakes was better than Pylint and PyChecker for 13 and 24 projects, respectively. PyChecker was better than Pylint for one project, by ratio. Standard deviations in mutation scores were sometimes high for Python: 0.17 for Pylint, 0.08 for pyflakes, and 0.01 for PyChecker.

For **RQ1**, there is a clear difference between tools. Pylint uniquely kills almost an order of magnitude more mutants than the next-best tool. There is a large overlap between Pylint and pyflakes, while PyChecker is both much less effective and

https://stackoverflow.com/questions/1428872/pylint-pychecker-or-pyflakes
https://www.reddit.com/r/Python/comments/ii3gm/experience_with_pylint_pychecker_pyflakes/
https://www.slant.co/versus/12630/12631/~pylint_vs_pyflakes
https://news.ycombinator.com/item?id=12748885
https://doughellmann.com/blog/2008/03/01/static-code-analizers-for-python/

Fig. 8: Discussions of Python static analysis tools.

doing something fairly different than the other tools. From the diagram, one might think that pyflakes acts, to some extent, as a less verbose "subset" of Pylint, in that most mutants detected by pyflakes are also detected by Pylint (however, the nearly 3,500 killed mutants unique to pyflakes suggest it is a useful tool, perhaps most useful after problems also reported by Pylint are fixed). PyChecker performs poorly in part because it crashed (due to changes in Python since the last update to the tool) for 19 of the 25 projects; however, it also performed poorly on programs where it worked.

The mutant ratios **RQ2** show that pyflakes is more competitive than is obvious from raw kill comparisons. The mutant ratio is almost three times as good as for Pylint! That is, some of Pylint's advantage may be due to general verbosity, even once stylistic warnings are turned off. Combining the results from raw kills and ratios, a strategy of using pyflakes as a quick check for problems, then using both pyflakes and Pylint for more in-depth analysis makes sense. Whether to use both pyflakes and Pylint in CI is a rquestion of tolerance for handling false positives, but it is clear that the "price" of additional warnings from Pylint is (1) high but (2) not without ereturn on investment (in terms of additional real finds). PyChecker is probably too outdated, and sometimes too verbose when it works, to be useful.

For **RQ3**, there were again no academic comparisons we could find. However, opinions on the web were quite common (see Figure 8, which lists ones we examined). It is hard to summarize the overall opinion here, since it ranges considerably. There is probably general agreement that PyChecker is old and maybe less useful, but also terse and sometimes helpful. Pylint is the most recommended tool, and the general complaint that it is too picky was somewhat mitigated in our results by turning off warnings that are clearly purely stylistic. Pyflakes is also well liked, and is considered much less verbose than Pylint; this was definitely reflected in our results, where Pyflakes underperformed in raw kills, but had a much better mutant ratio than Pylint. For **RQ4**, Pylint performed significantly better on clean projects. Performance for the other two tools was essentially unchanged.

For Python, all but three of the mutants we examined for **RQ 5** (a random sample of 100 mutants with a kill difference) involved code changes we agreed were definitely buggy, and would cause incorrect behavior if executed (the exceptions were deletions of code with no effect on state, e.g., strings as comments). The large majority involved statement deletions, detected via 1) unused variables/arguments, 2) instances lacking a member field, or 3) undefined variables. Interestingly, this seems to be an engine issue more than a detector issue,

as pyflakes and Pylint both basically support all of these kinds of checks. In some cases the two tools both detected a problem (but PyChecker did not) but differed as to *which* variable was not used or defined, again suggesting an engine rather than detection rule difference. Pylint's better performance was mostly, in the sample, due to detecting more of these issues, though it also was the only tool in the sample that detected arithmetic operation changes, due as far as we could determine to constant index changes. Pylint also detected a few mutants no other tool did, due to the presence of unreachable code. In one case, it also noticed a protected member access via a change in constant index, a surprisingly complex problem to find, in our opinion. PyChecker was never the sole detecting tool for any mutants in our sample.

D. Threats to Validity

The primary threat to validity in terms of generalization is that we only examined nine static analysis tools, and our analyses were restricted to 100 smart contracts, 15 Java projects, and 25 Python projects. Because it is hard to identify a ground truth to compare with (the motivation for our approach), we cannot be certain that our rankings of tools are correct even for these tools and this code. However, where there are existing discussions of the tools, our results seem to agree with these, but add substantial detail.

We used the Universal Mutator [27], [28], which aggressively produces large numbers of mutants, but does not target any particular software defect patterns, to generate all mutants. There is no room in the paper to present the exact set of projects analyzed, but we have provided an (anonymized) github repository containing raw results for inspection by reviewers, or further analysis by other researchers (https://github.com/mutantsforstaticanalysis/rawdata).

IV. RELATED WORK

The goal of "analysing the program analyser" [9] is intuitively attractive. The irony of using mostly ad-hoc, manual methods to test and understand static analysis tools is apparent; however, the fundamentally incomplete and heuristic nature of such tools makes this a challenge similar to testing machine learning algorithms [30]; most tools will not produce "the right answer" all the time, as a result of both algorithmic constraints and basic engineering trade-offs. While comparisons of static analysis tools [47], [18], [43], [19], [44] have appeared in the literature for years, these generally involved large human effort and resulting smaller scale, did not make a strong effort to address false positives, or restricted analysis to, e.g., a known defects set [32], [17]. Defect sets are vulnerable to tools intentionally overfitting/gaming the benchmark; our approach makes it easy to compare tools on "fresh" code to avoid this risk. Compared to well-known studies of Java tools [32], [47] our approach used a larger set of subject programs (1.8 MLOC total vs. 170-350KLOC) and thousands of *detected* faults, vs. e.g., about 500 known defects [32]. Results not using defect sets are even more limited in that humans can only realistically examine a few dozens of each type of warning [47].

Cuoq et al. [14] proposed the generation of random programs (*á la* Csmith [56]) to test analysis tools aiming for soundness, in limited circumstances, but noted that naïve differential testing of analysis tools was not possible. This paper proposes a non-naïve differential comparison based on the observation that the ability to detect program mutants offers an automatable comparison basis. We essentially adopt the approach of the large body of work on using mutants in software testing [34], [15], [7], [24], [25], [45], [33], [2], but re-define killing a mutation for a static analysis context.

Klinger et al. propose a different approach to differential testing of analysis tools [36]. Their approach is in some ways similar to ours, in that it takes as input a set of seed programs, and compares results across new versions generated from seeds. The primary differences are that their seed programs must be warning-free (which greatly limits the set of input programs available), and that the new versions are based on adding new assertions only. We allow arbitrarily buggy seed programs, and can, due to the any-language nature of the mutation generator we use, operate even in new languages. On the other hand, their approach can identify precision issues, while we offer no real help with false positives (in theory, you could apply their majority-vote method to mutants only a few tools flag, but mutants *are* usually faults. in contrast to their introduction of checks that may be guaranteed to pass). Most importantly, however, their approach *only applies to tools that check assertions*, rather than the much more common case of tools that identify bad code patterns.

V. CONCLUSIONS AND FUTURE WORK

In this paper, we showed that program mutants can be used as a proxy for real faults, to compare (and motivate improvements to) static analysis tools. Mutants are attractive in that a large body of work supports the claim that at least 60-70% of mutants are fault-inducing. This allows us to assume *detected* mutants are faulty, and escape the ground-truth/false positive problem that makes comparing static analysis tools so labor-intensive. We evaluated 9 popular static analysis tools, for Solidity smart contracts, Java, and Python, and offer advice to users of these tools. Our mutation results strongly confirm the wisdom of using multiple tools; with the exception of one Java tool, all tools we investigated uniquely detected over 1,000 mutants, and for Java and Python there were no mutants detected by all tools. For Solidity, academic research evaluations of the tools generally agreed strongly with our conclusions, but lacked the detail mutant analysis contributed. We were also able to use our methods, plus a novel mutant prioritization scheme, to identify and implement three useful new detectors for the open source Slither smart contract analyzer, the best-performing of the tools.

As future work, we would like to further validate our approach and improve our admittedly *ad hoc* mutant distance metric. Allowing user feedback [37], [30], or applying metric learning methods [38] (particularly unsupervised learning [48], [52]) are the most obvious and interesting possibilities.

REFERENCES

[1] Solidity 0.4.24 introduction to smart contracts. https://solidity.readthedocs.io/en/v0.4.24/introduction-to-smart-contracts.html.

[2] Iftekhar Ahmed, Carlos Jensen, Alex Groce, and Paul E. McKenney. Applying mutation analysis on kernel test suites: an experience report. In *International Workshop on Mutation Analysis*, pages 110–115, March 2017.

[3] James H. Andrews, L. C. Briand, and Y. Labiche. Is mutation an appropriate tool for testing experiments? In *International Conference on Software Engineering*, pages 402–411, 2005.

[4] Andreas M Antonopoulos and Gavin Wood. *Mastering Ethereum: building smart contracts and DApps.* O'Reilly Media, 2018.

[5] Nicola Atzei, Massimo Bartoletti, and Tiziana Cimoli. A survey of attacks on Ethereum smart contracts SoK. In *International Conference on Principles of Security and Trust*, pages 164–186, 2017.

[6] N. Ayewah, W. Pugh, D. Hovemeyer, J. D. Morgenthaler, and J. Penix. Using static analysis to find bugs. *IEEE Software*, 25(5):22–29, Sep. 2008.

[7] Timothy A Budd, Richard A DeMillo, Richard J Lipton, and Frederick G Sayward. Theoretical and empirical studies on using program mutation to test the functional correctness of programs. In *Principles of Programming Languages*, pages 220–233. ACM, 1980.

[8] Vitalik Buterin. Ethereum: A next-generation smart contract and decentralized application platform. https://github.com/ethereum/wiki/wiki/White-Paper, 2013.

[9] Cristian Cadar and Alastair F Donaldson. Analysing the program analyser. In *Proceedings of the 38th International Conference on Software Engineering Companion*, pages 765–768, 2016.

[10] T. T. Chekam, M. Papadakis, Y. Le Traon, and M. Harman. An empirical study on mutation, statement and branch coverage fault revelation that avoids the unreliable clean program assumption. In *2017 IEEE/ACM 39th International Conference on Software Engineering (ICSE)*, pages 597–608, May 2017.

[11] Yang Chen, Alex Groce, Chaoqiang Zhang, Weng-Keen Wong, Xiaoli Fern, Eric Eide, and John Regehr. Taming compiler fuzzers. In *Programming Language Design and Implementation*, pages 197–208, 2013.

[12] Henry Coles. Pit mutation testing: Mutators. http://pitest.org/quickstart/mutators.

[13] ConsenSys. Mythril: a security analysis tool for ethereum smart contracts. https://github.com/ConsenSys/mythril-classic, 2017.

[14] Pascal Cuoq, Benjamin Monate, Anne Pacalet, Virgile Prevosto, John Regehr, Boris Yakobowski, and Xuejun Yang. Testing static analyzers with randomly generated programs. In *NASA Formal Methods Symposium*, pages 120–125. Springer, 2012.

[15] Richard A DeMillo, Richard J Lipton, and Frederick G Sayward. Hints on test data selection: Help for the practicing programmer. *Computer*, 11(4):34–41, 1978.

[16] Dino Distefano, Manuel Fähndrich, Francesco Logozzo, and Peter W. O'Hearn. Scaling static analyses at facebook. *Commun. ACM*, 62(8):62–70, July 2019.

[17] Lisa Nguyen Quang Do, Michael Eichberg, and Eric Bodden. Toward an automated benchmark management system. In *Proceedings of the 5th ACM SIGPLAN International Workshop on State Of the Art in Program Analysis*, pages 13–17, 2016.

[18] Thomas Durieux, João F. Ferreira, Rui Abreu, and Pedro Cruz. Empirical review of automated analysis tools on 47,587 ethereum smart contracts. In *International Conference on Software Engineering*, 2020. Available as arXiv preprint at https://arxiv.org/abs/1910.10601.

[19] Josselin Feist, Gustavo Grieco, and Alex Groce. Slither: a static analyzer for solidity. In *International Workshop on Emerging Trends in Software Engineering for Blockchain*, pages 8–15, 2019.

[20] Omitted for blinding. Discussion of Java static analysis tools. Slack communication, January 21, 2020.

[21] Teofilo F. Gonzalez. Clustering to minimize the maximum intercluster distance. *Theoretical Computer Science*, 38:293–306, 1985.

[22] Rahul Gopinath, Amin Alipour, Iftekhar Ahmed, Carlos Jensen, and Alex Groce. How hard does mutation analysis have to be, anyway? In *2015 IEEE 26th International Symposium on Software Reliability Engineering (ISSRE)*, pages 216–227. IEEE, 2015.

[23] Rahul Gopinath, Carlos Jensen, and Alex Groce. Mutations: How close are they to real faults? In *International Symposium on Software Reliability Engineering*, pages 189–200, 2014.

[24] Alex Groce, Iftekhar Ahmed, Carlos Jensen, and Paul E McKenney. How verified is my code? falsification-driven verification. In *Automated Software Engineering (ASE), 2015 30th IEEE/ACM International Conference on*, pages 737–748. IEEE, 2015.

[25] Alex Groce, Iftekhar Ahmed, Carlos Jensen, Paul E McKenney, and Josie Holmes. How verified (or tested) is my code? falsification-driven verification and testing. *Automated Software Engineering Journal*, 25(4):917–960, 2018.

[26] Alex Groce, Josselin Feist, Gustavo Grieco, and Michael Colburn. What are the actual flaws in important smart contracts (and how can we find them)? In *International Conference on Financial Cryptography and Data Security*, 2020. Accepted for publication.

[27] Alex Groce, Josie Holmes, Darko Marinov, August Shi, and Lingming Zhang. Regexp based tool for mutating generic source code across numerous languages. https://github.com/agroce/universalmutator.

[28] Alex Groce, Josie Holmes, Darko Marinov, August Shi, and Lingming Zhang. An extensible, regular-expression-based tool for multi-language mutant generation. In *International Conference on Software Engineering: Companion Proceeedings*, pages 25–28, 2018.

[29] Alex Groce, Gerard Holzmann, and Rajeev Joshi. Randomized differential testing as a prelude to formal verification. In *International Conference on Software Engineering*, pages 621–631, 2007.

[30] Alex Groce, Todd Kulesza, Chaoqiang Zhang, Shalini Shamasunder, Margaret Burnett, Weng-Keen Wong, Simone Stumpf, Shubhomoy Das, Amber Shinsel, Forrest Bice, and Kevin McIntosh. You are the only possible oracle: Effective test selection for end users of interactive machine learning systems. *IEEE Transactions on Software Engineering*, 40(3):307–323, March 2014.

[31] B. J. M. Grün, D. Schuler, and A. Zeller. The impact of equivalent mutants. In *2009 International Conference on Software Testing, Verification, and Validation Workshops*, pages 192–199, April 2009.

[32] Andrew Habib and Michael Pradel. How many of all bugs do we find? a study of static bug detectors. In *Proceedings of the 33rd ACM/IEEE International Conference on Automated Software Engineering*, ASE 2018, page 317–328, New York, NY, USA, 2018. Association for Computing Machinery.

[33] Goran Petrović Marko Ivanković, Bob Kurtz, Paul Ammann, and René Just. An industrial application of mutation testing: Lessons, challenges, and research directions. In *Proceedings of the International Workshop on Mutation Analysis (Mutation). IEEE Press, Piscataway, NJ, USA*, pages 47–53, 2018.

[34] Yue Jia and Mark Harman. An analysis and survey of the development of mutation testing. *IEEE Transactions on Software Engineering*, 37(5):649–678, 2011.

[35] René Just, Darioush Jalali, and Michael D Ernst. Defects4J: A database of existing faults to enable controlled testing studies for java programs. In *Proceedings of the 2014 International Symposium on Software Testing and Analysis*, pages 437–440. ACM, 2014.

[36] Christian Klinger, Maria Christakis, and Valentin Wüstholz. Differentially testing soundness and precision of program analyzers. In *Proceedings of the 28th ACM SIGSOFT International Symposium on Software Testing and Analysis*, pages 239–250, 2019.

[37] Todd Kulesza, Margaret M. Burnett, Simone Stumpf, Weng-Keen Wong, Shubhomoy Das, Alex Groce, Amber Shinsel, Forrest Bice, and Kevin McIntosh. Where are my intelligent assistant's mistakes? A systematic testing approach. In *End-User Development - Third International Symposium, IS-EUD 2011, Torre Canne (BR), Italy, June 7-10, 2011. Proceedings*, pages 171–186, 2011.

[38] Brian Kulis. Metric learning: A survey. *Foundations & Trends in Machine Learning*, 5(4):287–364, 2012.

[39] Duc Le, Mohammad Amin Alipour, Rahul Gopinath, and Alex Groce. MuCheck: An extensible tool for mutation testing of Haskell programs. In *Proceedings of the 2014 International Symposium on Software Testing and Analysis*, pages 429–432. ACM, 2014.

[40] Vladimir I. Levenshtein. Binary codes capable of correcting deletions, insertions, and reversals. *Soviet Physics Doklady*, 10:707–710, 1966.

[41] William McKeeman. Differential testing for software. *Digital Technical Journal of Digital Equipment Corporation*, 10(1):100–107, 1998.

[42] Mike Papadakis, Yue Jia, Mark Harman, and Yves Le Traon. Trivial compiler equivalence: A large scale empirical study of a simple fast and effective equivalent mutant detection technique. In *International Conference on Software Engineering*, 2015.

[43] Reza M. Parizi, Ali Dehghantanha, Kim-Kwang Raymond Choo, and Amritraj Singh. Empirical vulnerability analysis of automated smart

contracts security testing on blockchains. In *Proceedings of the 28th Annual International Conference on Computer Science and Software Engineering*, CASCON '18, page 103–113, USA, 2018. IBM Corp.

[44] Ivan Pashchenko, Stanislav Dashevskyi, and Fabio Massacci. Delta-bench: differential benchmark for static analysis security testing tools. In *2017 ACM/IEEE International Symposium on Empirical Software Engineering and Measurement (ESEM)*, pages 163–168. IEEE, 2017.

[45] Goran Petrović and Marko Ivanković. State of mutation testing at google. In *Proceedings of the 40th International Conference on Software Engineering: Software Engineering in Practice*, ICSE-SEIP '18, pages 163–171, New York, NY, USA, 2018. ACM.

[46] Phil Daian . Analysis of the dao exploit. http://hackingdistributed.com/2016/06/18/analysis-of-the-dao-exploit/, June 18, 2016 (acceded on Jan 10, 2019).

[47] Nick Rutar, Christian B. Almazan, and Jeffrey S. Foster. A comparison of bug finding tools for Java. In *Proceedings of the 15th International Symposium on Software Reliability Engineering*, ISSRE '04, page 245–256, USA, 2004. IEEE Computer Society.

[48] Bernhard Schölkopf, Alexander Smola, and Klaus-Robert Müller. Nonlinear component analysis as a kernel eigenvalue problem. *Neural computation*, 10(5):1299–1319, 1998.

[49] Ben H Smith and Laurie Williams. Should software testers use mutation analysis to augment a test set? *Journal of Systems and Software*, 82(11):1819–1832, 2009.

[50] SpankChain. We got spanked: What we know so far. https://medium.com/spankchain/we-got-spanked-what-we-know-so-far-d5ed3a0f38fe, Oct 8, 2018 (acceded on Jan 10, 2019).

[51] Sergei Tikhomirov, Ekaterina Voskresenskaya, Ivan Ivanitskiy, Ramil Takhaviev, Evgeny Marchenko, and Yaroslav Alexandrov. Smartcheck: Static analysis of ethereum smart contracts. In *International Workshop on Emerging Trends in Software Engineering for Blockchain*, pages 9–16, 2018.

[52] Michael E Tipping and Christopher M Bishop. Probabilistic principal component analysis. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 61(3):611–622, 1999.

[53] David A. Tomassi, Naji Dmeiri, Yichen Wang, Antara Bhowmick, Yen-Chuan Liu, Premkumar T. Devanbu, Bogdan Vasilescu, and Cindy Rubio-González. BugSwarm: mining and continuously growing a dataset of reproducible failures and fixes. In *International Conference on Software Engineering*, pages 339–349. IEEE / ACM, 2019.

[54] Petar Tsankov, Andrei Dan, Dana Drachsler-Cohen, Arthur Gervais, Florian Bünzli, and Martin Vechev. Securify: Practical security analysis of smart contracts. In *ACM SIGSAC Conference on Computer and Communications Security*, pages 67–82, 2018.

[55] Gavin Wood. Ethereum: a secure decentralised generalised transaction ledger. http://gavwood.com/paper.pdf, 2014.

[56] Xuejun Yang, Yang Chen, Eric Eide, and John Regehr. Finding and understanding bugs in C compilers. In *Programming Language Design and Implementation*, pages 283–294, 2011.